

# Lecture 1. Course Introduction and Practicalities

## IL2233 Embedded Intelligence

Prof. Zhonghai Lu

KTH Royal Institute of Technology

March 19, 2024

# Outline

1. What is the course about?
2. Course design principles
3. Course structure, content and organization
4. Course examination
5. Practicalities
6. Demo: Sine wave prediction

# What is the course about?

- This is a new course (from 2022), developed to meet the educational needs of developing and deploying various intelligent data analysis functions in sensor-based embedded systems, Internet of Things (IoT), and Cyber-Physical Systems (CPS).
  - IL2230 Hardware Architectures for Deep Learning from 2019.
- The course aims to provide the students with **essential theoretical methods and practical skills**, which are needed to develop, assess, and deploy intelligent functionalities in smart electronics and embedded systems. In particular, **dependability** and **sustainability** are considered in the course.

# What to focus on?

- The course covers *selected statistical learning/machine learning/deep learning methods* for realizing relevant functionalities (e.g. feature exaction, time-series forecasting, clustering, and anomaly detection, etc.) desired in embedded systems.
- The course focuses more on *applied time-series analysis techniques*, touching upon both statistical and deep learning approaches, and *their application in embedded systems*.
  - Intelligent embedded systems often deal with time-series data from various sensors such as temperature/humidity, pressure, camera (image sensor), IMU (Inertial Measurement unit), lidar, radar, EEG (Electroencephalogram), ECG (Electrocardiogram) etc., much often than time-independent data.
  - Time-series analysis has its own set of theory and tools, in particular, from the statistical learning domain.

# Course design principles

- Balance between theory and practice, but more on practice and practical skills.
- In terms of breadth and depth, the course is broader, covering a wide range of knowledge.  
In some areas, the depth is limited, leaving space for own exploration.
- Focus on understanding of basic concepts and principles rather than complicated application scenarios, so as to establish a good foundation for further investigations on your own.
- Hands-on coding examples.

# Course structure

The course is designed to have 16 teaching & learning activities.

- 10 Lectures organized as 3 modules
- 3 Labs (Lab1, Lab2, Lab3 + Lab0 individual preparation for tool installation)
- 2 Seminars

- 1 Project

A common application of time-series analysis for embedded system: Anomaly detection

# Course content

- The course content is very carefully selected, balanced, and designed.
  - It has gone through many times of revision.
  - In aware of workload, relevance, significance, etc.
- The Canvas course room <https://canvas.kth.se/courses/46239>
- Let's go to the canvas course room to see
  - What are the modules, lectures?
  - What are the labs?
  - What are the seminars?
  - What is the project?

# Activity organization

- 3 Labs
  - Group work, 2 students per group.
  - Lab completion: 1 technical report per group. Results check per group
  - Grouping: Automatic/Random grouping in Canvas
- 2 Seminars
  - Group work, 3-4 students per group.
  - Seminar: Being **presenter** and **opponent**
    - Each group makes presentation slides with everyone's contribution.
    - During presentation, each one presents one part of the presentation.
    - In Q & A, each group also acts as opponent group to another group.
  - Grouping: Automatic/Random grouping in Canvas
- 1 Project
  - **Individual work**, Individual completion.
  - There are two Project sessions for you to do the project tasks. Similar to Lab sessions, TAs are available for assistance.



# Examination

- No written examination but the **Final Project Workshop** works as Oral Examination.
  - It is compulsory, requiring physical presence.
  - In the Project Workshop, you make slides, orally present your project individually, and answer questions.
- Three examination moments
  - Lab
  - Seminar
  - Project
- You need to pass all the 3 moments in order to pass the course
- Course grading: Pass, Fail
  - Allowing self-exploration without worrying about grades

# Background survey

- This link is for students to answer the survey.  
<https://forms.gle/YmwoT59Ry7QRKjmv5>

# Programming environment

- Programming language: Python, C/C++
- Basic Python libraries
  - numpy, scipy
  - matplotlib
  - seaborn: Statistical data visualization <https://seaborn.pydata.org/>
  - pandas
- Statistical/Machine learning toolbox
  - Statistical learning library: statsmodels, pmdarima
  - Machine learning: scikit-learn
  - Deep learning library: keras, tensorflow

# Python: A high-level programming language

- Python is an interpreted language, which can save you considerable time during program development because no compilation and linking is necessary.
- Python allows you to split your program into modules that can be reused in other Python programs.

# Python: A high-level programming language

- Python is simple to use, and enables programs to be written compactly and readably.
- Programs written in Python are typically much shorter than equivalent C, C++, or Java programs, for several reasons:
  - the high-level data types allow you to express complex operations in a single statement;
  - statement grouping is done by indentation instead of beginning and ending brackets;
  - no variable or argument declarations are necessary.
- The Python tutorial.

`https://docs.python.org/3/tutorial/index.html`

# Packages

- Packages are a way of structuring Python's module namespace by using “dotted module names”.
- For example, the module name A.B designates a submodule named B in a package named A.

# Numpy

- Numpy is a core library for **Scientific computing** in Python.
- Numpy provides a high-performance multidimensional array and basic tools to compute with and manipulate arrays, including mathematical, logical, shape manipulation, sorting, selecting, I/O, discrete Fourier transforms, basic linear algebra, basic statistical operations, random simulation and much more.
- Numpy quick start.

`https://numpy.org/doc/stable/user/quickstart.html`

# Scipy

- SciPy, a scientific library for Python is an open source, BSD-licensed library for mathematics, science and engineering.
- The Scipy library depends on NumPy, and provides a large number of functions that operate on numpy arrays.
- Scipy user guide: an introductory tutorial, which covers the fundamentals of SciPy and describes how to deal with its various modules.

<https://docs.scipy.org/doc/scipy/tutorial/index.html>



# Matplotlib

- Matplotlib is a Python 2D plotting library. which produces **publication quality** figures in a variety of hardcopy formats and interactive environments across platforms.
- Matplotlib can be used in Python scripts, the Python and IPython shells, the Jupyter notebook, web application servers, and graphical user interface toolkits.
- Matplotlib  
`https://matplotlib.org/index.html`

# Matplotlib

- One can generate plots, histograms, power spectra, bar charts, error charts, scatter plots, etc., with just a few lines of code.
- For simple plotting the **pyplot** module provides a MATLAB-like interface, particularly when combined with IPython.
- For advanced use, you have full control of line styles, font properties, axes properties, etc, via an object oriented interface or via a set of functions familiar to MATLAB users.
- Pyplot tutorial: An introduction to the pyplot interface.

`https:`

`//matplotlib.org/stable/tutorials/introductory/pyplot.html`

# Pandas

- Pandas is a software library written in Python for **data manipulation** and **analysis**.
- In particular, it offers data structures and operations for manipulating numerical tables and time series.
- Pandas user guide and 10 minutes to pandas  
[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/](https://pandas.pydata.org/pandas-docs/stable/user_guide/)

# Scikit-learn (sklearn)

- A **machine-learning** library or package in Python.
- It features various classification, regression and clustering algorithms including support-vector machines, random forests, k-means etc.
  - Simple and efficient tools for predictive data analysis
  - Built on NumPy, SciPy, and matplotlib
  - Open source, commercially usable - BSD license
- Scikit-learn tutorial

`https://scikit-learn.org/stable/tutorial/index.html`

# Keras: Deep Learning Library for Theano and TensorFlow

- Keras is a **deep learning** API written in Python, running on top of the machine learning platform TensorFlow or Theano. Or, Keras is a high-level API of the tensorflow/Theano platform.
- It was developed with a focus on enabling fast experimentation from training to inference.
- Keras website.  
`https://keras.io/about/`
- Keras means **horn** in Greek.

# Keras: Deep Learning Library for Theano and TensorFlow

- The core data structures of Keras are **layers** and **models**.
- The simplest type of model is the Sequential model, a linear stack of layers.
- For more complex architectures, you should use the Keras functional API, which allows to build arbitrary graphs of layers, or write models entirely from scratch via sub-classing.

# Statsmodels

- statsmodels is a Python module that provides classes and functions for the estimation of many different **statistical models**, as well as for conducting **statistical tests**, and **statistical data exploration**.
- Website  
`https://www.statsmodels.org/stable/index.html`
- Tutorial  
`https://www.statsmodels.org/stable/gettingstarted.html`

# General installation guide

- Download and install **Anaconda**.
  - This gives you Python along with many built-in libraries.
- Create and activate a **virtual environment**, e.g. IL2233  
`https://uoa-eresearch.github.io/eresearch-cookbook/recipe/2014/11/20/conda/`
  - This is an additional environment to the default environment.
  - This is important for keeping an environment to use a specific version of a library.
- If a package is not present, install it using either **conda** or **pip**
- Use the virtual environment
  - Activate the virtual environment
  - De-activate the virtual environment



# Installation of sklearn

- Use conda  
conda install -c anaconda scikit-learn
- Use pip  
pip install scikit-learn  
pip install -U scikit-learn  
pip install scikit-learn --upgrade  
-U or --upgrade: Upgrade all packages to the newest available version.

<https://scikit-learn.org/stable/install.html>

# Installation of keras

Keras comes packaged with TensorFlow as tensorflow.keras.

To start using Keras, simply install TensorFlow: `pip install tensorflow`

- Use conda  
`conda install -c anaconda keras`
- Use pip  
`pip install keras`
- If you already have TensorFlow and Keras installed, they can be updated and then verified by:  
`pip install -U tensorflow`  
`python -m pip show tensorflow`

<https://keras.io/about/>

# Installation of statsmodels

- Installing statsmodels

- The easiest way to install statsmodels is to install it as part of the Anaconda distribution, a cross-platform distribution for data analysis and scientific computing. This is the recommended installation method for most users.

- Python support

- statsmodels supports Python 3.8, 3.9, and 3.10.

- Anaconda

- statsmodels is available through conda provided by Anaconda.
- The latest release can be installed using:  
`conda install -c conda-forge statsmodels`

- PyPI (pip)

- To obtain the latest released version of statsmodels using pip:  
`pip install statsmodel`

<https://www.statsmodels.org/devel/install.html>

# Validate your installation

- Show if you have a package, which version
  - `pip show scikit-learn`
  - `pip show keras`
  - `pip show tensorflow`
  - `pip show statsmodels`
- List all installed Python packages
  - `pip list`
- Show Python/module version
  - `python --version`
  - `pip --version`
  - `anaconda --version`

# Demo: Sine wave prediction (1)

Task: Train a neural network (NN) using a synthetic dataset from a noisy  $y(t)=\sin(t)$  function, and then predict the values of the sine function  $y$  given values of  $t$ .

Implementation: Generic learning/inference steps:

- 1 Data generation: Generate data by a noisy  $\sin()$  function:  $y(t) = \sin(t) + c \cdot \epsilon(t)$ , where  $\epsilon(t)$  is Gaussian noise with mean 0 and variance 1, and  $c$  coefficient for modulating the noise.
- 2 Data pre-processing: Split the data set into a training set and a test set, e.g. 80% for training and 20% for testing.

## Demo: Sine wave prediction (2)

- 1 Model definition: Define a NN model, e.g., a MLP (multi-layer perceptron) model including its hyper-parameters (number of layers and number of neurons per layer, activation function etc.) and optimization parameters (optimizer (e.g. Adam: Adaptive moment estimation), loss function).
- 2 Model training: Train the NN to settle its weights and biases, optimizing towards the loss function.
- 3 Model inference: Use the trained model to do prediction given an argument value.
- 4 Model evaluation: Calculate accuracy metrics of the trained model.
- 5 Result visualization. Visualize results.

`http://localhost:8889/notebooks/OneDrive%20-%20KTH/IPython/IL2233VT22/Lec1_intro/sine_wave_prediction.ipynb`

(The link is only used for in-classroom demo, not accessible by students.)

# Homework

- Installation of Python and the packages.
  - Use Anaconda
  - Create a virtual work environment
- Go through the Python tutorial  
`https://docs.python.org/3/tutorial/index.html`
- Go through other tutorials such as numpy, scipy, matplotlib etc.
- If you are all good, try to implement the sine wave prediction task.
- **Thinking question:** What if we treat the sine wave as a **time-series** signal, and how to predict, more precisely, forecast its next value, based on its present and previous values?